

# Ai\_lab06\_TahaAbbas\_P200119

March 16, 2023

## Lab Task 06 Multi layer Perceptron

### *Predicting Who Survived the Titanic Using Multilayer Perceptron (MLP)*

The goal of this lab task is to build a multilayer perceptron (MLP) classifier that can predict whether a passenger survived the sinking of the Titanic or not.

1. Import the necessary Python libraries, such as pandas, numpy, and sklearn.
2. Load the Titanic dataset into a pandas DataFrame using `pandas.read_csv()`.
3. Preprocess the data by converting categorical features into numerical ones, filling in missing values, and scaling the numerical features using `sklearn.preprocessing`.
4. Split the dataset into training and test sets using `sklearn.model_selection.train_test_split()`.
5. Build an MLP classifier using `sklearn.neural_network.MLPClassifier()` and train it on the training data.
6. Evaluate the performance of the MLP classifier on the test data using metrics such as accuracy, precision, recall, and F1-score.
7. Fine-tune the MLP classifier by adjusting its hyperparameters, such as the number of hidden layers, and the number of neurons per layer.
8. Evaluate the performance of the fine-tuned MLP classifier on the test data and compare it to the initial model.
9. Discuss the results and insights gained from the experiment, and identify potential areas for further improvement.

```
[1]: #imports
import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Visualization imports
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[ ]:
```

```
[2]: df = pd.read_csv('titanic_train.csv')
df.head()
```

```
[2]:
```

	passenger_id	pclass	name \
0	1216	3	Smyth, Miss. Julia
1	699	3	Cacic, Mr. Luka
2	1267	3	Van Impe, Mrs. Jean Baptiste (Rosalie Paula Go...
3	449	2	Hocking, Mrs. Elizabeth (Eliza Needs)
4	576	2	Veal, Mr. James

	sex	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body \
0	female	NaN	0	0	335432	7.7333	NaN	Q	13	NaN
1	male	38.0	0	0	315089	8.6625	NaN	S	NaN	NaN
2	female	30.0	1	1	345773	24.1500	NaN	S	NaN	NaN
3	female	54.0	1	3	29105	23.0000	NaN	S	4	NaN
4	male	40.0	0	0	28221	13.0000	NaN	S	NaN	NaN

	home.dest	survived
0	NaN	1
1	Croatia	0
2	NaN	0
3	Cornwall / Akron, OH	1
4	Barre, Co Washington, VT	0

```
[3]: df.shape
```

```
[3]: (850, 15)
```

```
[4]: df.columns
```

```
[4]: Index(['passenger_id', 'pclass', 'name', 'sex', 'age', 'sibsp', 'parch',
          'ticket', 'fare', 'cabin', 'embarked', 'boat', 'body', 'home.dest',
          'survived'],
          dtype='object')
```

```
[5]: df.describe()
```

```
[5]:
```

	passenger_id	pclass	age	sibsp	parch \
count	850.000000	850.00000	676.000000	850.000000	850.000000
mean	662.816471	2.32000	29.519847	0.522353	0.382353
std	380.751936	0.83853	14.562243	1.112132	0.879511
min	1.000000	1.00000	0.166700	0.000000	0.000000
25%	332.250000	2.00000	20.000000	0.000000	0.000000
50%	676.500000	3.00000	28.000000	0.000000	0.000000
75%	992.250000	3.00000	37.000000	1.000000	0.000000
max	1307.000000	3.00000	80.000000	8.000000	9.000000

	fare	body	survived
count	849.000000	73.000000	850.000000
mean	34.012701	165.821918	0.368235
std	53.705779	99.068487	0.482610
min	0.000000	4.000000	0.000000
25%	7.895800	75.000000	0.000000
50%	14.108300	166.000000	0.000000
75%	31.000000	260.000000	1.000000
max	512.329200	328.000000	1.000000

```
[6]: df.isnull().sum()
```

```
[6]: passenger_id      0
     pclass            0
     name              0
     sex               0
     age              174
     sibsp             0
     parch             0
     ticket            0
     fare              1
     cabin            659
     embarked          1
     boat             542
     body             777
     home.dest         386
     survived          0
     dtype: int64
```

```
[7]: df.survived.value_counts()
```

```
[7]: 0    537
     1    313
     Name: survived, dtype: int64
```

```
[8]: df.embarked.value_counts()
```

```
[8]: S    589
     C    176
     Q     84
     Name: embarked, dtype: int64
```

```
[9]: df.pclass.value_counts()
```

```
[9]: 3    478
     1    206
     2    166
```

Name: pclass, dtype: int64

```
[10]: label_encoder = LabelEncoder()

df['sex'] = label_encoder.fit_transform(df['sex'])
df['embarked'] = label_encoder.fit_transform(df['embarked'].astype(str))
df.head()
```

```
[10]:
```

	passenger_id	pclass	name \
0	1216	3	Smyth, Miss. Julia
1	699	3	Cacic, Mr. Luka
2	1267	3	Van Impe, Mrs. Jean Baptiste (Rosalie Paula Go...
3	449	2	Hocking, Mrs. Elizabeth (Eliza Needs)
4	576	2	Veal, Mr. James

	sex	age	sibsp	parch	ticket	fare	cabin	embarked	boat	body \
0	0	NaN	0	0	335432	7.7333	NaN	1	13	NaN
1	1	38.0	0	0	315089	8.6625	NaN	2	NaN	NaN
2	0	30.0	1	1	345773	24.1500	NaN	2	NaN	NaN
3	0	54.0	1	3	29105	23.0000	NaN	2	4	NaN
4	1	40.0	0	0	28221	13.0000	NaN	2	NaN	NaN

	home.dest	survived
0	NaN	1
1	Croatia	0
2	NaN	0
3	Cornwall / Akron, OH	1
4	Barre, Co Washington, VT	0

```
[11]: # Select features and target variable
X = df[['pclass', 'sex', 'age', 'sibsp', 'parch', 'fare', 'embarked']]
y = df['survived']
```

```
[12]: X.head(10)
```

```
[12]:
```

	pclass	sex	age	sibsp	parch	fare	embarked
0	3	0	NaN	0	0	7.7333	1
1	3	1	38.0	0	0	8.6625	2
2	3	0	30.0	1	1	24.1500	2
3	2	0	54.0	1	3	23.0000	2
4	2	1	40.0	0	0	13.0000	2
5	3	1	28.0	0	0	22.5250	2
6	3	1	19.0	0	0	0.0000	2
7	2	0	30.0	0	0	13.0000	2
8	3	0	22.0	0	0	7.7750	2
9	3	0	21.0	1	0	9.8250	2

```
[13]: X.isnull().sum()
```

```
[13]: pclass      0
      sex        0
      age      174
      sibsp      0
      parch      0
      fare       1
      embarked   0
      dtype: int64
```

```
[14]: #df["age"] = df["age"].fillna(df["age"].mean())
```

```
[ ]: X.fillna(X.mean(), inplace=True)
```

```
[16]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 850 entries, 0 to 849
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   pclass      850 non-null    int64
 1   sex         850 non-null    int64
 2   age         850 non-null    float64
 3   sibsp       850 non-null    int64
 4   parch       850 non-null    int64
 5   fare        850 non-null    float64
 6   embarked    850 non-null    int64
dtypes: float64(2), int64(5)
memory usage: 46.6 KB
```

```
[ ]: scaler = StandardScaler()
```

```
column = ['age', 'sibsp', 'parch', 'fare']
X[column] = scaler.fit_transform(X[column])
```

```
[18]: X.head()
```

```
[18]:   pclass  sex      age      sibsp      parch      fare  embarked
0       3    0  2.737723e-16 -0.469963 -0.434989 -0.489898         1
1       3    1  6.534810e-01 -0.469963 -0.434989 -0.472576         2
2       3    0  3.700060e-02  0.429741  0.702675 -0.183860         2
3       2    0  1.886442e+00  0.429741  2.978005 -0.205298         2
4       2    1  8.076011e-01 -0.469963 -0.434989 -0.391717         2
```

```
[19]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=10)
```

```
[20]: X_train.head(3)
```

```
[20]:
```

	pclass	sex	age	sibsp	parch	fare	embarked
476	2	0	-1.966561	1.329444	0.702675	0.092973	2
411	3	1	-0.502420	-0.469963	-0.434989	-0.502636	2
377	1	0	-0.810660	-0.469963	1.840340	-0.144091	2

```
[21]: X_test.head(3)
```

```
[21]:
```

	pclass	sex	age	sibsp	parch	fare	embarked
780	3	1	2.737723e-16	-0.469963	-0.434989	-0.471644	2
233	3	1	2.737723e-16	-0.469963	-0.434989	0.419128	2
68	3	1	-4.005945e-02	-0.469963	-0.434989	-0.457275	2

```
[22]: y_train.value_counts()
```

```
[22]: 0    425
      1    255
      Name: survived, dtype: int64
```

```
[23]: y_test.value_counts()
```

```
[23]: 0    112
      1     58
      Name: survived, dtype: int64
```

```
[24]: mlp = MLPClassifier(solver='lbfgs', alpha=1e-5, hidden_layer_sizes=(8,6),
↳max_iter=1000, random_state=1)
      mlp.fit(X_train, y_train)
      y_pred = mlp.predict(X_test)
```

### Before Fine tuning

```
[25]: print('Accuracy score:', round(accuracy_score(y_test, y_pred),2))
      print('Precision score:', round(precision_score(y_test, y_pred),2))
      print('Recall score:', round(recall_score(y_test, y_pred),2))
      print('F1 score:', round( f1_score(y_test, y_pred),2))
```

```
Accuracy score: 0.78
Precision score: 0.75
Recall score: 0.52
F1 score: 0.61
```

```
[ ]: mlp = MLPClassifier(solver='lbfgs', alpha=1e-5, hidden_layer_sizes=(8,8),
    ↪max_iter=1500, random_state=1)
mlp.fit(X_train, y_train)
y_pred = mlp.predict(X_test)
```

After tuning


```
[27]: print('Accuracy score:', round(accuracy_score(y_test, y_pred),2))
print('Precision score:', round(precision_score(y_test, y_pred),2))
print('Recall score:', round(recall_score(y_test, y_pred),2))
print('F1 score:', round( f1_score(y_test, y_pred),2))
```

Accuracy score: 0.79  
Precision score: 0.78  
Recall score: 0.55  
F1 score: 0.65


```
[28]: Xlim = df[['pclass', 'sex', 'age', 'sibsp', 'parch', 'fare', 'embarked',
    ↪'survived']]
```

```
[29]: sns.barplot(data=Xlim, x="sex", y="survived")
```



```
[29]: <Axes: xlabel='sex', ylabel='survived'>
```







```
[30]: for column in Xlim.columns:
        fig, ax = plt.subplots()
        ax.bar(df[column].value_counts().index, df[column].value_counts().
        ↪values)
        ax.set_xlabel(column)
        ax.set_ylabel('Survived')
        plt.show()
```














*We can observe from the results that the fine-tuned model performs marginally better than the baseline MLP model in terms of accuracy, precision, recall, and F1-score. This shows that modifying the hyperparameters can enhance model performance, and that more feature engineering and data preparation may do the same.*

[ ]: